INTRODUCTION TO REGRESSION ANALYSIS

Gianluca Campanella

Senior ML Researcher, Longshot Systems Ltd

LEARNING OBJECTIVES

‣ Define data modeling and simple linear regression

‣ Build a linear regression model using a dataset that meets the linearity assumption using the sci-kit learn library

‣ Understand and identify multicollinearity in a multiple regression.

INTRODUCTION TO REGRESSION ANALYSIS

PRE-WORK

PRE-WORK REVIEW

‣ Effectively show correlations between an independent variable x and a dependent variable y

‣ Be familiar with the get_dummies function in pandas

‣ Understand the difference between vectors, matrices, Series, and DataFrames

‣ Be able to interpret p-values and confidence intervals

INTRODUCTION TO REGRESSION ANALYSIS

WHERE ARE WE IN THE DATA SCIENCE WORKFLOW?

‣ Data has been **acquired** and **parsed**

‣ Today we'll **refine** the data and **build** models

‣ We'll also use plots to **represent** the results

SIMPLE LINEAR REGRESSION

SIMPLE LINEAR REGRESSION

‣ Explanation of a continuous variable given a series of independent variables

‣ The simplest version is just a line of best fit:
y = mx + b

‣ Explain the relationship between **x** and **y** using the starting point **b** and the power in explanation **m**.

SIMPLE LINEAR REGRESSION

‣ However, linear regression uses linear algebra to explain the relationship between *multiple* x's and y.

‣ The more sophisticated version: $y = \beta * X$ (+ error)

‣ Explain the relationship between the matrix **X** and a dependent vector **y** using a y-intercept $\boldsymbol{\beta_0}$ and the relative coefficients $\boldsymbol{\beta_{1, 2, \ldots}}$.

SIMPLE LINEAR REGRESSION

‣ Linear regression works **best** when:

  ‣ The data is normally distributed (but doesn't have to be)

  ‣ X's significantly explain y (have low p-values)

  ‣ X's are independent of each other (low multicollinearity)

  ‣ Resulting values pass linear assumption (depends upon problem)

‣ If data is not normally distributed, we could introduce *bias*

REGRESSING AND NORMAL DISTRIBUTIONS

DEMO: REGRESSING AND NORMAL DISTRIBUTIONS

‣ Follow along with your starter code notebook while I walk through these examples

‣ The first plot shows a relationship between two values, though not a linear solution

‣ Note that lmplot() returns a straight line plot

‣ However, we can transform the data, both log-log distributions to get a linear solution

USING SEABORN TO GENERATE SIMPLE LINEAR MODEL PLOTS

ACTIVITY: GENERATE SINGLE VARIABLE LINEAR MODEL PLOTS

DIRECTIONS (15 minutes)

1. Update and complete the code in the starter notebook to use **lmplot** and display correlations between body weight and two dependent variables: **sleep_rem** and **awake**.

EXERCISE

DELIVERABLE

Two plots

SIMPLE REGRESSION ANALYSIS IN SKLEARN

SIMPLE LINEAR REGRESSION ANALYSIS IN SKLEARN

‣ Sklearn defines models as *objects* (in the OOP sense)

‣ You can use the following principles:

  ‣ All sklearn modeling classes take a similar form

  ‣ All estimators take a matrix $\mathbf{X}$, either sparse or dense

  ‣ Supervised estimators also take a vector $\mathbf{y}$ (the response)

  ‣ Estimators can be customized through setting the appropriate parameters

CLASSES AND OBJECTS IN OBJECT ORIENTED PROGRAMMING

‣ **Classes** are an abstraction for a complex set of ideas, e.g. *human*.

‣ Specific **instances** of classes can be created as **objects**.
  ‣ *john_smith = human()*

‣ Objects have **properties**. These are attributes or other information.
  ‣ *john_smith.age*
  ‣ *john_smith.gender*

‣ Object have **methods**. These are procedures associated with a class/object.
  ‣ *john_smith.breathe()*
  ‣ *john_smith.walk()*

SIMPLE LINEAR REGRESSION ANALYSIS IN SKLEARN

‣ General format for sklearn model classes and methods

```python
# generate an instance of an estimator class
estimator = base_models.AnySKLearnObject()
# fit your data
estimator.fit(X, y)
# score it with the default scoring method (recommended to use the metrics module in the future)
estimator.score(X, y)
# predict a new set of data
estimator.predict(new_X)
# transform a new X if changes were made to the original X while fitting
estimator.transform(new_X)
```

‣ LinearRegression() doesn't have a transform function

‣ With this information, we can build a simple process for linear regression.

SIGNIFICANCE IS KEY

DEMO: SIGNIFICANCE IS KEY

‣ Follow along with your starter code notebook while I walk through these examples.

‣ What does the residual plot tell us?

‣ How can we use the linear assumption?

USING THE LINEAR REGRESSION OBJECT

ACTIVITY: USING THE LINEAR REGRESSION OBJECT

EXERCISE

DIRECTIONS (15 minutes)

1. With a partner, generate two more models using the log-transformed data to see how this transform changes the model's performance.
2. Use the code on the following slide to complete #1.

DELIVERABLE

Two new models

ACTIVITY: USING THE LINEAR REGRESSION OBJECT

EXERCISE

DIRECTIONS (15 minutes)

```python
X =
y =
loop = []
for boolean in loop:
    print 'y-intercept:', boolean
    lm =
linear_model.LinearRegression(fit_intercept=boolean)
    get_linear_model_metrics(X, y, lm)
    print
```

DELIVERABLE

Two new models

BASE LINEAR REGRESSION CLASSES

ACTIVITY: BASE LINEAR REGRESSION CLASSES

EXERCISE

DIRECTIONS (20 minutes)

1. Experiment with the model evaluation function we have (**get_linear_model_metrics**) with the following sklearn estimator classes.

   a. linear_model.Lasso()
   b. linear_model.Ridge()
   c. linear_model.ElasticNet()

Note: We'll cover these new regression techniques in a later class.

## DELIVERABLE

New models and evaluation metrics

MULTIPLE REGRESSION ANALYSIS

MULTIPLE REGRESSION ANALYSIS

‣ Simple linear regression with one variable can explain some variance, but using multiple variables can be much more powerful.

‣ We want our multiple variables to be mostly independent to avoid multicollinearity.

‣ Multicollinearity, when two or more variables in a regression are highly correlated, can cause problems with the model.

BIKE DATA EXAMPLE

‣ We can look at a correlation matrix of our bike data.

‣ Even if adding correlated variables to the model improves overall variance, it can introduce problems when explaining the output of your model.

‣ What happens if we use a second variable that isn't highly correlated with temperature?

MULTICOLLINEARITY WITH DUMMY VARIABLES

ACTIVITY: MULTICOLLINEARITY WITH DUMMY VARIABLES

EXERCISE

DIRECTIONS (15 minutes)

1. Load the bike data.
2. Run through the code on the following slide.
3. What happens to the coefficients when you include all weather situations instead of just including all except one?

DELIVERABLE

Two models' output

ACTIVITY: MULTICOLLINEARITY WITH DUMMY VARIABLES

DIRECTIONS (15 minutes)

EXERCISE

```python
lm = linear_model.LinearRegression()
weather = pd.get_dummies(bike_data.weathersit)
get_linear_model_metrics(weather[[1, 2, 3, 4]], y, lm)
print
# drop the least significant, weather situation  = 4
get_linear_model_metrics(weather[[1, 2, 3]], y, lm)
```

DELIVERABLE

Two models' output

COMBINING FEATURES INTO A BETTER MODEL

ACTIVITY: COMBINING FEATURES INTO A BETTER MODEL

EXERCISE

DIRECTIONS (15 minutes)

1. With a partner, complete the code on the following slide.
2. Visualize the correlations of all the numerical features built into the dataset.
3. Add the three significant weather situations into our current model.
4. Find two more features that are not correlated with the current features, but could be strong indicators for predicting guest riders.

DELIVERABLE

Visualization of correlations, new models

ACTIVITY: COMBINING FEATURES INTO A BETTER MODEL

EXERCISE

DIRECTIONS (15 minutes)

```python
lm = linear_model.LinearRegression()
bikemodel_data = bike_data.join() # add in the three weather situations

cmap = sns.diverging_palette(220, 10, as_cmap=True)
correlations = # what are we getting the correlations of?
print correlations
print sns.heatmap(correlations, cmap=cmap)

columns_to_keep = [] #[which_variables?]
final_feature_set = bikemodel_data[columns_to_keep]

get_linear_model_metrics(final_feature_set, y, lm)
```

DELIVERABLE

Visualization of correlations, new models

BUILDING MODELS FOR OTHER Y VARIABLES

ACTIVITY: BUILDING MODELS FOR OTHER Y VARIABLES

EXERCISE

DIRECTIONS (25 minutes)

1. Build a new model using a new y variable: registered riders.
2. Pay attention to the following:
   a. the distribution of riders (should we rescale the data?)
   b. checking correlations between the variables and y variable
   c. choosing features to avoid multicollinearity
   d. model complexity vs. explanation of variance
   e. the linear assumption

BONUS

1. Which variables make sense to dummy?
2. What features might explain ridership but aren't included? Can you build these features with the included data and pandas?

DELIVERABLE

A new model and evaluation metrics

TOPIC REVIEW

CONCLUSION

‣ You should now be able to answer the following questions:

   ‣ What is simple linear regression?

   ‣ What makes multi-variable regressions more useful?

   ‣ What challenges do they introduce?

   ‣ How do you dummy a category variable?

   ‣ How do you avoid a singular matrix?

INTRODUCTION TO REGRESSION ANALYSIS

Q & A